

# CS747 - Assignment 4

Kalpesh Krishna  
140070017  
kalpeshk2011@gmail.com


## 1 Implementation & Results

All algorithms were implemented in Python without the use of any external library except `numpy` to generate random numbers and carry out matrix updates. The `weights` array has been offset by 1 to start indices from 1 (instead of 0) to match Baird's counterexample diagram. A random seed 0 has been used in all experiments.

### 1.1 Experiment 1


In this experiment, all non-terminal states are chosen in a round robin fashion followed by a TD(0) style update. The figure below shows the graph obtained. This is similar in shape to the graph obtained in chapter 8 of Sutton and Barto 1998.

**Explanation** - The reason for this anomalous behaviour lies in the updates of  $w[7]$ . Initially, all weights are 1, so  $\gamma V[6] < V[1..5]$ , which cause  $w[1..5], w[7]$  to decrease.  $w[6]$  almost never updates due to the 99% self loop. But  $w[1..5]$  decrease faster than  $w[7]$  and eventually  $\gamma V[6] > V[1..5]$ . This causes an increase in  $V(6)$ , until  $\gamma V[6] < V[1..5]$ , and this cycle continues. Since the value of  $V(6)$  has a stronger dependence on  $w[7]$  when compared to  $V[1..5]$ , continuous updates to state 1..5 in response to a strong positive or negative reward due to  $\gamma V[6]$  makes  $w[7]$  go to positive infinity, and hence  $w[1..5]$  go to positive infinity.  $w[6]$  on the other hand goes to negative infinity in an attempt to counteract the effect of  $w[7]$ . But the infrequent updates to  $w[6]$  slow this down.



## 1.2 Experiment 2

No significant difference was observed by changing the value of  $\lambda$ . Unlike the previous case,  $V[1...5]$  are updated less often.  $w[1...5]$  are updated exactly once for each non-zero update in  $w[6]$  in the TD(0) case. In the previous task,  $w[6]$  is non-zero updated roughly once in 600 iterations, whereas  $w[1...5]$  are updated 100 times each.



## 1.3 Experiment 3

I noticed that changing the initial set of weights did not affect the final value function achieved. However, some configurations were taking longer to converge, which I suspect is due to a constant learning rate. The final set of weights were different, they differed by a constant factor. However, the ratio of  $w[1...5] : w[6] : w[7]$  was always roughly  $1 : 4 : -2$ , which is expected behavior due to the choice of function approximation. Raw output has been added below.

```
./startmdp.sh 2 1000000 0 1 1 1 1 1 1 1
```

Weights

```
[ 0.27999309  0.28000883  0.28002641  0.28001831  0.2799872   1.11999228
 -0.55999853]
```

Value Function

```
-1.23541347308e-05 1.91397635615e-05 5.42867724099e-05 3.8093030442e-05 -2.4125545401e-05 -4.7
```

```
./startmdp.sh 2 1000000 0 3 1 10 1 10 1 2
```

Weights

```
[ 0.99940458  0.99880368  1.00166754  0.9984671   1.00126741  4.00009876
 -1.99999732]
```

Value Function

```
-0.0011881549616 -0.00238995587097 0.00333776830031 -0.00306312343532 0.00253750595343 0.00010
```

```
./startmdp.sh 2 1000000 0 100 1 0 1 4 80 3
```

Weights

```
[ 16.82400073  16.79381982  16.79406201  16.79204024  16.79566909
```

67.20004989 -33.60010428]

Value Function

0.0478971806153 -0.0124646358046 -0.0119802537092 -0.0160238019396 -0.0087660968456 -0.0001586

./startmdp.sh 2 1000000 0 1000 1 0 1 4 80 3

Weights

[ 53.04225399 52.73819236 52.74732618 52.72276639 52.74817208

211.19984362 -105.60095726]

Value Function

0.483550712958 -0.124572552079 -0.106304900622 -0.1554244781 -0.104613113732 -0.00207090789942

./startmdp.sh 2 1000000 0 1000 1 0 1 4 800 3

Weights

[ 168.24305459 167.93828292 167.94312627 167.92048556 167.95005989

672.00069251 -336.00111036]

Value Function

0.484998817983 -0.124544515748 -0.114857826734 -0.160139244123 -0.100990581897 -0.001528210843